

Chapter 15

Multiple Service Servers

Outline:

Motivation

Iterative, Connectionless Server.

Concurrent, Connection-Oriented Server.

Single-Process Server

Flexibility of Invoking Separate Programs

Multiservice, Multiprotocol Designs

Example Multiservice Server

BSD super server: inetd

Motivation:

Consolidation of servers reduces:

- the number of executing processes
- the total code required

Principle: use a separate socket (and port) for each service, but have one program handle them all.

Otherwise similar to the corresponding single service Server.

Iterative, Connectionless, Multiservice Server

- 1 Opens a set of UDP sockets.
- 2 Binds each one to a well-known port for one of the services being offered.
- 3 A table maps services to sockets (for each descriptor, the table records the address of a procedure to call for service).
- 4 Call select to wait for a datagram to arrive.
- 5 Examine the descriptor and call the appropriate service procedure.
- 6 Return to the select call.

Iterative is good for handling short, non-interactive services.

Iterative, Connection-Oriented, Multiservice Server

- 1 Opens one socket for each service and binds each one to the appropriate well-known port.
- 2 A table maps services to sockets (for each descriptor, the table records the address of a procedure)
- 3 Calls select to wait for an incoming connection request.
- 4 When a request arrives, calls accept to create a new socket for the incoming connection.
- 5 Uses the new socket to interact with the client, then closes the socket.
- 6 Return to the select call.

Concurrent, Connection-Oriented, Single-Process Multiservice Server

- 1 Opens one master socket for each service and binds each one to the appropriate well-known port.
- 2 A table maps services to sockets (for each descriptor, the table records the address of a procedure)
- 3 Calls select to wait for incoming connection requests master sockets or incoming data to read on slave sockets.
- 4a For connect requests, calls accept to create a new socket for the incoming connection.
- 4b For incoming data, calls the appropriate procedure to provide service.
- 5 Return to the select call.

Concurrent, Connection-Oriented, Multi-Process Multiservice Server

- 1 Opens one master socket for each service and binds each one to the appropriate well-known port.
- 2 A table maps services to sockets (for each descriptor, the table records the address of a procedure)
- 3 Calls select to wait for incoming connection requests.
- 4 Creates a slave processes for each new connection.
- 5 Each slave handles service for client.

Software Design

Break code into independent components using separately compiled programs.

Compile those modules which change and relink.

Multi-process Servers: `fork/exec` allow call of a separately compiled (and maintained program).

Example Multiservice Server

see Section 15.9 for code

Four services: echo, chargen, daytime, and time

```
struct service{
    char *sv_name;          /* ASCII name of service */
    char sv_useTCP;         /* 0=UDP 1=TCP          */
    int sv_sock;            /* socket descriptor   */
    int (*sv_func)(int);    /* function to call     */
}
```

svent is an array of 4 of these

Each entry is used to store the information about one of the services.

C note: sv_func stores the address of a function.

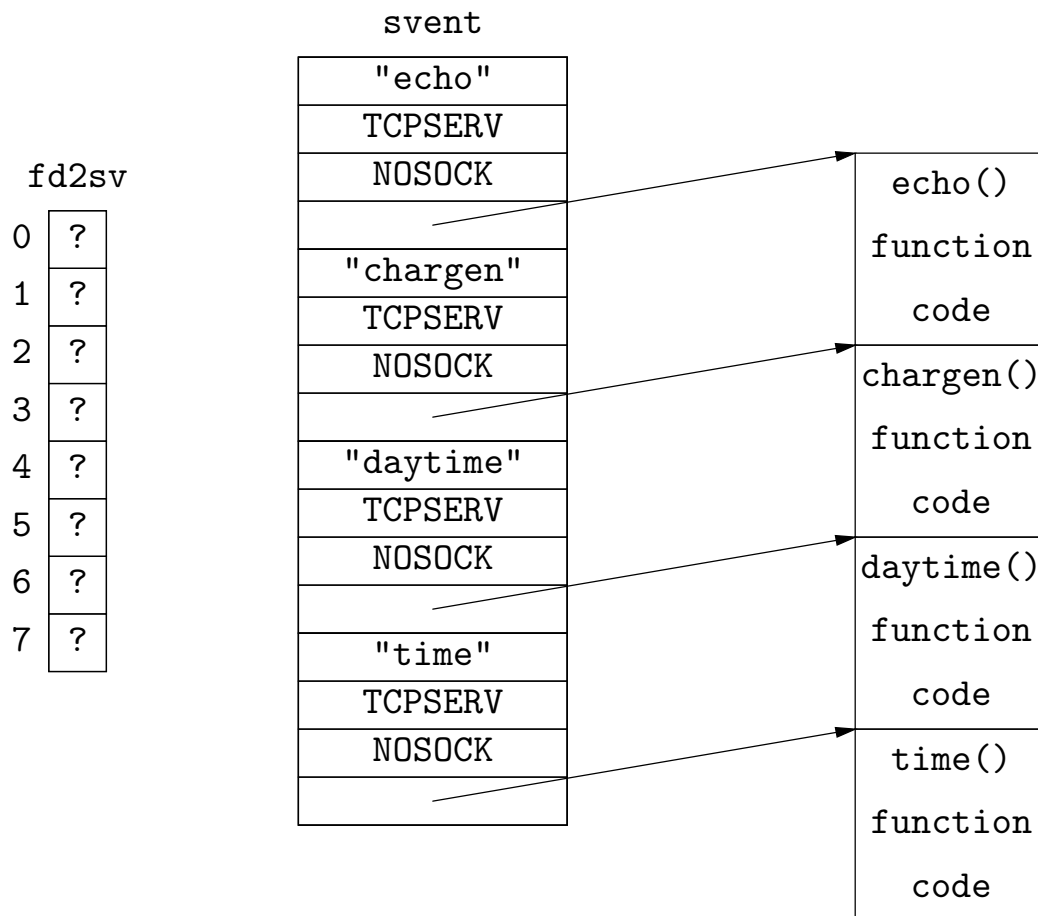
function calls are by address

svent[3]->sv_func(4) will call the function whose address is found in the 4th record (TCPtimed), and ask it to use the socket found in descriptor 4.

svent[3]->sv_sock is the master socket for the service

Array fd2sv to map from the descriptor back to an entry in svent (reverse map).

Structure Before socket allocation

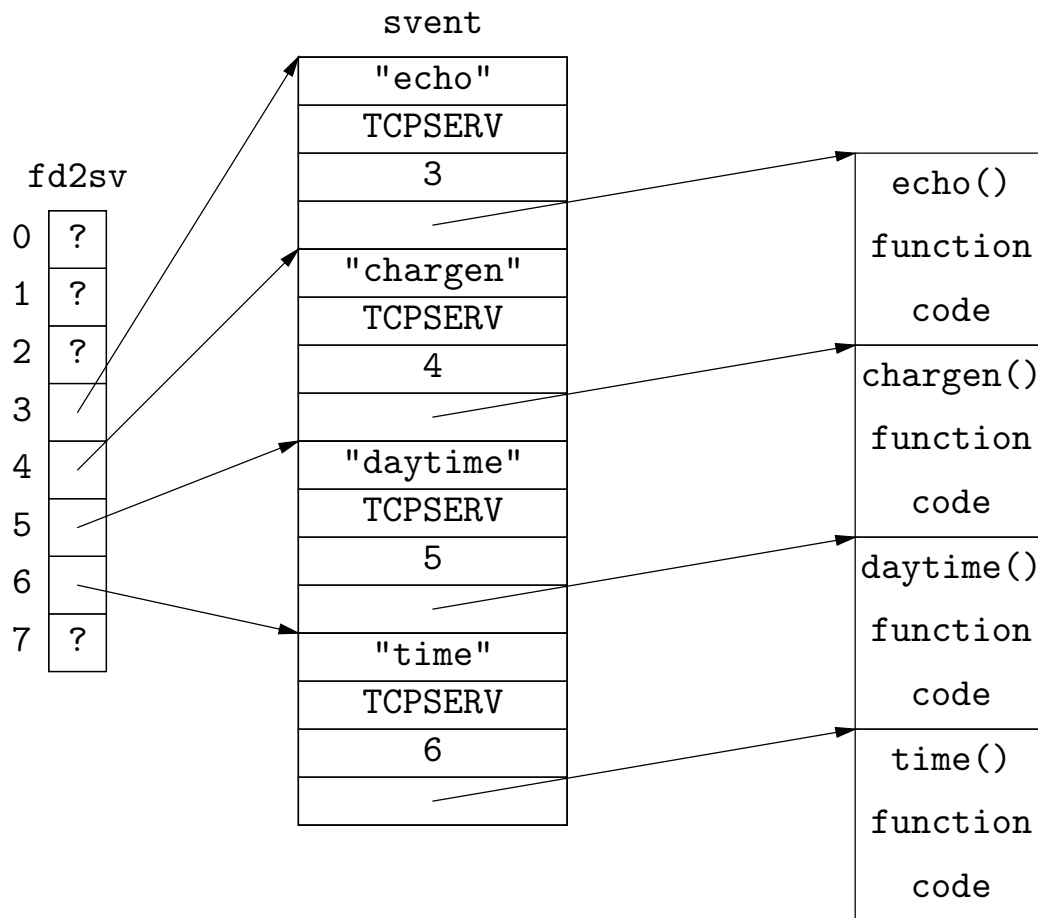


Strings are for human interface purposes and for use with the `passiveUDP` and `passiveTCP` functions.

No sockets are entered (yet)

0,1,2: are `stdin`, `stdout` and `stderr` and will be unused

Structure after socket allocation



Section 15.9 code details

```
#define  UDP_SERV  0
#define  TCP_SERV  1
#define  NOSOCK -1 /*an invalid socket descriptor*/
struct service {
    char    *sv_name;
    char    sv_useTCP;
    int     sv_sock;
    int     (*sv_func)(int);
};
struct service svent[] = {
    {"echo",TCP_SERV,NOSOCK,TCPecho },
    {"chargen",TCP_SERV,NOSOCK,TCPchargend },
    {"daytime",TCP_SERV,NOSOCK,TCPdaytimed },
    {"time",TCP_SERV,NOSOCK,TCPtimed },
    {0, 0, 0, 0 }
};
```

The defines set up values to go into the structure.

The structure is declared and initialized to have everything but the descriptor (socket) number.

```

int main(int argc, char *argv[]) {
    struct service *psv,    /*service table pointer*/
        *fd2sv[NOFILE];    /*map fd to service pointer*/
    int  fd, nfds;
    fd_set  afd, rfd;
    nfds = getdtablesize();
    FD_ZERO(&afd);
    for (psv = &svent[0]; psv->sv_name; ++psv) {
        if (psv->sv_useTCP)
            psv->sv_sock = passiveTCP(psv->sv_name, QLEN);
        else
            psv->sv_sock = passiveUDP(psv->sv_name);
        fd2sv[psv->sv_sock] = psv;
        FD_SET(psv->sv_sock, &afd);
    }
}

```

for: go through the array of services and fill in the socket numbers.

a 0 for the name indicates the end of the array

Open either a TCP master socket or
a UDP socket for each service.

Also, fill in the reverse lookup array fd2sv

```
while (1) {  
    memcpy(&rfd, &afd, sizeof(rfd));  
    if (select(nfds, &rfd, NULL, NULL, NULL) < 0) {  
        if (errno == EINTR) continue;  
        errexit(...);  
    }  
}
```

while(1) the main service loop,

memcpy...select: wait for a connect/message on any of the ports we have open.

EINTR: Falling out the select because the program has received a SIGCHLD signal is handled by going back around the while loop to the select. This is necessary since we have fork and a reaper.

```

for (fd=0; fd<nfds; ++fd)
    if (FD_ISSET(fd, &rfdsets)) {
        psv = fd2sv[fd];
        if (psv->sv_useTCP)
            doTCP(psv);
        else
            psv->sv_func(psv->sv_sock);
    }
}

```

for process each connect/message.

This is the real service inside the service loop.

Use the reverse map (`fd2sv`) to find the service entry information

If the service entry indicates this is a UDP function, just call the service function passing the socket.

If the service entry indicates this is a TCP function, call `doTCP` because TCP needs additional setup before calling the service procedure

```

void doTCP(struct service *psv) {
    struct sockaddr_in fsin;
    int  alen, ssock;
    alen = sizeof(fsin);
    ssock=accept(psv->sv_sock,
        (struct sockaddr *)&fsin, &alen);
    if (ssock < 0) errexit(...);
    switch (fork()) {
    case 0:  break;
    case -1: errexit(...);
    default: (void) close(ssock);
        return;    /* parent */
    }
    exit(psv->sv_func(ssock)); /* child */
}

```

This is the special TCP setup.

do an accept to set up the slave socket

fork a child that calls the service function

when the service function returns, the child exits

return takes the parent back to the main program, which
for-loops to the next rfd's needing service.

The service function must be correctly programmed.

For TCP, the service function is passed the slave socket on which it must do any necessary reads and writes.

The `accept` has been done, the `close` is not required because of the `exit`

For UDP, the service function must do a `recvfrom` so has an address to use for the `sendto`.

When done, the service function should return; which will trigger the `exit`.

inetd

BSD UNIX super server

Supplies many small TCP/IP services

Services can be added to (deleted from) inetd using a configuration file

/etc/inetd.conf the config file

Service name—such as telnet

Socket type—dgram/stream

Service type—udp/tcp

wait/nowait—single process/multiprocess (fork)

user—privilege level (actually user name)

service procedure—name of the executable file, internal indicates a tiny service like time provided directly by inetd.

display name—what to show to ps

Only one server is running waiting for many things.

Specific servers are not started unless needed.