

Internet Protocol Version 6

The version 6 of the IP protocol is designed to replace the current version 4 of the protocol.

Notes:

IP Addresses change from 32 bits to 128 bits.

There are backward compatibility features:

- 1) old addresses translate to new addresses: 96 bits of 0.
- 2) Dual stack host can send and receive both.
- 3) Dual stack router can forward both.
- 4) Dual protocol routers can translate packets from/to IPv4 to/from IPv6

The IPv6 definitions are already available in the `.h` files.

Almost all Linux machines are dual stack.

Activate by assigning IP4 and IP6 addresses.

IPv6 structures (in.h)

The address is 128 bits (was `sin_addr`).

In both versions it is stored inside a struct.

In version 4 it was a long so the nesting didn't matter.

In version 6 placing the array in a struct avoids the complications arising from the fact that arrays are represented by pointers.

```
struct in6_addr {  
    u_int8_t  s6_addr[16];  
}
```

In the version6 address structure, the family and port are the same as before, just newer type names for the same sizes. The family will be `AF_INET6`.

The flow information variable is new, exactly how it is to be used is not fully defined.

The address portion uses the new structure.

```
struct sockaddr_in6 {  
    u_char          sin6_family;  
    u_int16m_t      sin6_port;  
    u_int32m_t      sin6_flowinfo;  
    struct in6_addr sin6_addr;  
};
```

There are version 6 values for the constants. (`in.h`)

Initializing the address structure is similar to version 4.

```
struct sockaddr_in6 sin6;
/* Server */
sin6.sin6_family = AF_INET6;
sin6.sin6_flowinfo = 0;
sin6.sin6_port = htons(23);
sin6.sin6_addr = in6addr_any;
/* Client with server on local host */
sin6.sin6_family = AF_INET6;
sin6.sin6_flowinfo = 0;
sin6.sin6_port = htons(23);
sin6.sin6_addr = in6addr_loopback;
```

The API calls take version 6 parameters

```
s = socket(AF_INET6, SOCK_STREAM, IPPROTO_IPV6);
if (bind(s, (struct sockaddr *)&sin6, sizeof(sin6)) < 0)
{
    errexit(...);
}
```

The socket options have version 6 equivalents.

There are a few additional options.

```
setsockopt(fd, IPPROTO_IPV6, IPV6_HOPLIMIT,
    &one, sizeof(one));
```

Names to Numbers

The system administrator can set up a machines Domain Name Service so that `gethostbyname` returns IPv6 addresses.

(Admin Detail) `/etc/resolv.conf` add the line
`options inet6`

Same call:

```
hp = gethostbyname(...);
```

Returned structure:

`hp->h_length` will be 16 (was 4)

Change the memcpy code to use a `sin6`

```
memcpy(&sin6.sin6_addr, hp->h_addr, hp->h_length);
```

Problem: support both IP4 and IP6 addresses

Solution: use `getaddrinfo` instead.

getaddrinfo

Designed as a generalization of (and improvement on) `gethostbyname/addr` and `getservicebyname/port`.

The following structure is used by `getaddrinfo`. It is used both to specify the nature of the request (IP4, IP6, ...) and to return the answer.

```
struct addrinfo {  
    int ai_flags;  
    int ai_family;  
    int ai_socktype;  
    int ai_protocol;  
    int ai_addrlen;  
    struct sockaddr *ai_addr;  
    char *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

`ai_family`: `AF_INET` or `AF_INET6`

`ai_socktype`: `SOCK_STREAM` or `SOCK_DGRAM`

`ai_protocol`: The protocol to use.

`ai_flags`: control what is returned

`ai_addr`: the address returned

`ai_canonname`: the fully qualified host name

`ai_next`: because sizes can vary a linked list is returned.

```
int getaddrinfo(const char *host,  
                const char *service,  
                const struct addrinfo *hints,  
                struct addrinfo **res);
```

host: The name or IP number of a host. IP number can be given in IP4 dotted notation, IP6 byte notation or IP6 mixed notation (see hint). If this is `NULL`, then the appropriate loop back address will be returned.

service: The name (eg. "telnet") or number (eg. "7654") of a service. If this is `NULL` the port number in the result will not be initialized.

hints: Specifies the type of lookup should be done.

res: The information being returned (result).

RETURN value: 0 if no error, otherwise the value is an error code that can be used with `gai_strerror` to get a string describing the error.

Details

hints:

`ai_family` specifies whether a IP4 or IP6 address should be returned. A value of `AF_UNSPEC` returns either or both (linked list).

`ai_socktype` UDP/TCP, 0 means retrieve all.

`ai_flags` what should be returned

`AI_CANONNAME` return the canonical name

`AI_ADDRCONFIG` used with `unspec`, return only addresses this host is configured to use. (Don't return an IPv6 if we have no IPv6 interfaces configured.)

res:

This is created by the `getaddrinfo` call. Consequently it needs to be released using a `freeaddrinfo` call if you are doing a bunch of get addrs. The `ai_addr` field is an entire address structure, including the socket type, port number and address; for the appropriate family (IP4/6).

A linked of list structures is returned if more than one address matches the hints.

strings to/from addresses

New calls are needed to turn a string into an address or an address into a string

String to address:

`inet_pton` Turn a string into an address

```
inet_pton(AF_INET6, "::134.139.248.19",&a.sin6_addr);
```

```
inet_pton(AF_INET6, "::8f8b:f813",&a.sin6_addr);
```

The `::` indicates 0's.

`a` must be a `sockaddr_in6`

Address to string:

`inet_ntop` Turn an address into a string

```
inet_ntop(AF_INET6,info_ptr->ai_addr,b,sizeof(b));
```

`b` should be a character array of size at least

`INET6_ADDRSTRLEN`.

Primitive Client

```
int main() {
    int s;
    char message[80];
    struct sockaddr_in6  srv_addr;
    s = socket(PF_INET6, SOCK_STREAM, 0);
    memset(&srv_addr, 0, sizeof(srv_addr));
    inet_pton(AF_INET6, "::134.139.248.19",
        &srv_addr.sin6_addr);
    srv_addr.sin6_family = AF_INET6;
    srv_addr.sin6_port = htons(7654);
    connect(s, (struct sockaddr *) &srv_addr,
        sizeof(srv_addr));
    strcpy(message, "Client speaks");
    write(s, message, 80);
    read(s, message, 80);
    close(s);
    return 1;
}
```

Primitive Client Notes

Changes:

1) lots of "6"s

2) used `inet_pton`

Primitive Client 2

```
int main() {
    int s;
    char message[80];
    char addr_buf[INET6_ADDRSTRLEN]; /* for ntop */
    struct sockaddr_in6  srv_addr;
    struct addrinfo hint, *infoPtr;
    int result;
    hint.ai_family = AF_INET6;
    hint.ai_socktype = SOCK_STREAM;
    hint.ai_protocol = 0; /* any protocol */
    hint.ai_flags = AI_CANONNAME | AI_ADDRCONFIG;
    result = getaddrinfo("puma",NULL,&hint,&infoPtr);
    memcpy(&srv_addr,infoPtr->ai_addr,
        infoPtr->ai_addrlen);
    srv_addr.sin6_port = htons(7654);
    s = socket(PF_INET6, SOCK_STREAM, 0);
    connect(s, (struct sockaddr *) &srv_addr,
        sizeof(srv_addr));
    strcpy(message,"Client speaks");
    write(s, message, 80);
    read(s, message, 80);
    close(s);
    return 1;
}
```

Primitive Client 2 Notes

Changes:

- 1) lots of "6"s
- 2) used `getaddrinfo`
- 3) mem copied the entire structure, not just the address part.

Could have:

```
getaddrinfo("puma", "7654", &hint, &info_ptr);
```

instead of doing it separately

Primitive server

```
int main(){
    int master, client, len; char message[80];
    struct sockaddr_in6 my_addr, his_addr;
    /* get a socket allocated */
    master = socket(PF_INET6, SOCK_STREAM, 0);
    /* bind to the well-known port on our machine */
    memset(&my_addr, 0, sizeof(my_addr));
    my_addr.sin6_family = AF_INET6;
    my_addr.sin6_flowinfo = 0;
    my_addr.sin6_addr = in6addr_any;
    my_addr.sin6_port = htons(7654);
    bind(master, (struct sockaddr *) &my_addr,
        sizeof(my_addr));
    listen(master, 5);
    len=sizeof(his_addr);
    /* get the connection to the client */
    client = accept(master,
        (struct sockaddr *) &his_addr, &len);
    /* get the message from the client */
    read(client,message,80);
    strcpy(message,"Server replies");
    write(client, message, 80); /* send reply */
    close(client);
    close(master);
    return 1;
}
```

Primitive Server Notes

Changes:

1) lots of "6"s

2) in6addr_any